

Teertha Reddy Maddika

+1 (716) 335 2001 • Buffalo, NY • Portfolio • teerthareddy.m@gmail.com • LinkedIn • GitHub

SUMMARY

AI/ML Engineer with **2+ years of experience** designing and deploying **production-grade machine learning systems**. Proficient in building **real-time data pipelines** using **AWS services** such as S3, Glue, Lambda, and deploying **optimized models** with SageMaker, TensorRT. Skilled in **data engineering, workflow orchestration, and cloud-native deployment** with strong foundations in **computer vision** and **NLP**.

TECHNICAL SKILLS

Programming Languages: Python, PySpark, SQL, Java, C++.

ML Frameworks: PyTorch, TensorFlow, HuggingFace, Keras.

Cloud & Data Platforms: AWS (S3, Glue, Lambda, EMR, Redshift, SageMaker), Snowflake, Kafka, Airflow, Docker

Tools: MLflow, dbt, Power BI, Git, Spark, Jupyter.

Soft Skills: Communication, Agile teamwork, Problem-solving.

WORK EXPERIENCE

Machine Learning Engineer

Hyderabad, TS, India

Nexora AI Solutions

Jun 2021 – Jul 2023

- Built and deployed scalable **ML pipelines** using **AWS Glue, EMR, Lambda, and SageMaker**, reducing model latency from hours to **milliseconds**.
- Developed **production-ready models** in **PyTorch** and **Scikit-learn**, achieving **>93% accuracy**.
- Designed **real-time streaming** solutions with **Kinesis + Lambda**, processing **IoT data** at **<5s latency**.
- **Streamlined model ops** using **Airflow (retraining) + MLflow (tracking)**, reducing overhead by **45%.**
- Migrated **ETL** to **distributed processing** using **PySpark** and optimized **Spark clusters**, reducing costs by **40%.**
- Built **40+ dbt models** for **feature transformation** and **SCD2 history tracking**, improving reproducibility.
- Deployed **data validation checks** using **Great Expectations**, enhancing trust in **ML inputs**.
- Collaborated across teams to **productionize models** with **CI/CD**, monitored through **Prometheus + Grafana**.

PROJECTS

AI-Powered Feedback Analyzer for Education Platforms

Jan 2025

Python, HuggingFace, Streamlit, FastAPI

- Built **NLP pipeline** using **DistilBERT** to detect sentiment and inclusivity in education platform feedback
- Achieved **92% F1-score** on fine-tuned classifier and deployed dashboard via **Streamlit + FastAPI**.

Cloud-Native Fraud Detection System

May 2024

AWS Kinesis, SageMaker, PySpark, Redshift

- Designed **real-time fraud detection pipeline** using **Kinesis Firehose + Lambda**, handling **1M+ events/day**.
- Used **SageMaker Pipelines** for full **ML lifecycle**; results integrated with **Redshift + Power BI**.
- Reduced **false positives by 28%** through calibrated **threshold tuning and post-processing**.

Smart Resource Allocator (Capstone)

Dec 2023

TensorFlow, Kubernetes, Flask, Prometheus

- Built **ML model** to predict workload spikes and **scale pods dynamically** with **K8s HPA + Prometheus metrics**.
- Achieved **85% prediction accuracy**, deployed using **Flask REST API** for live inference.

CERTIFICATIONS

- **AWS Certified Machine Learning – Specialty** — 2024
- **Google Professional ML Engineer** — 2024
- **NVIDIA Deep Learning Certificate** — 2024

EDUCATION

University at Buffalo – SUNY

Buffalo, NY, USA.

MS in Computer Science (AI/ML Track) CGPA: 3.5

Jan 2024 – May 2025.